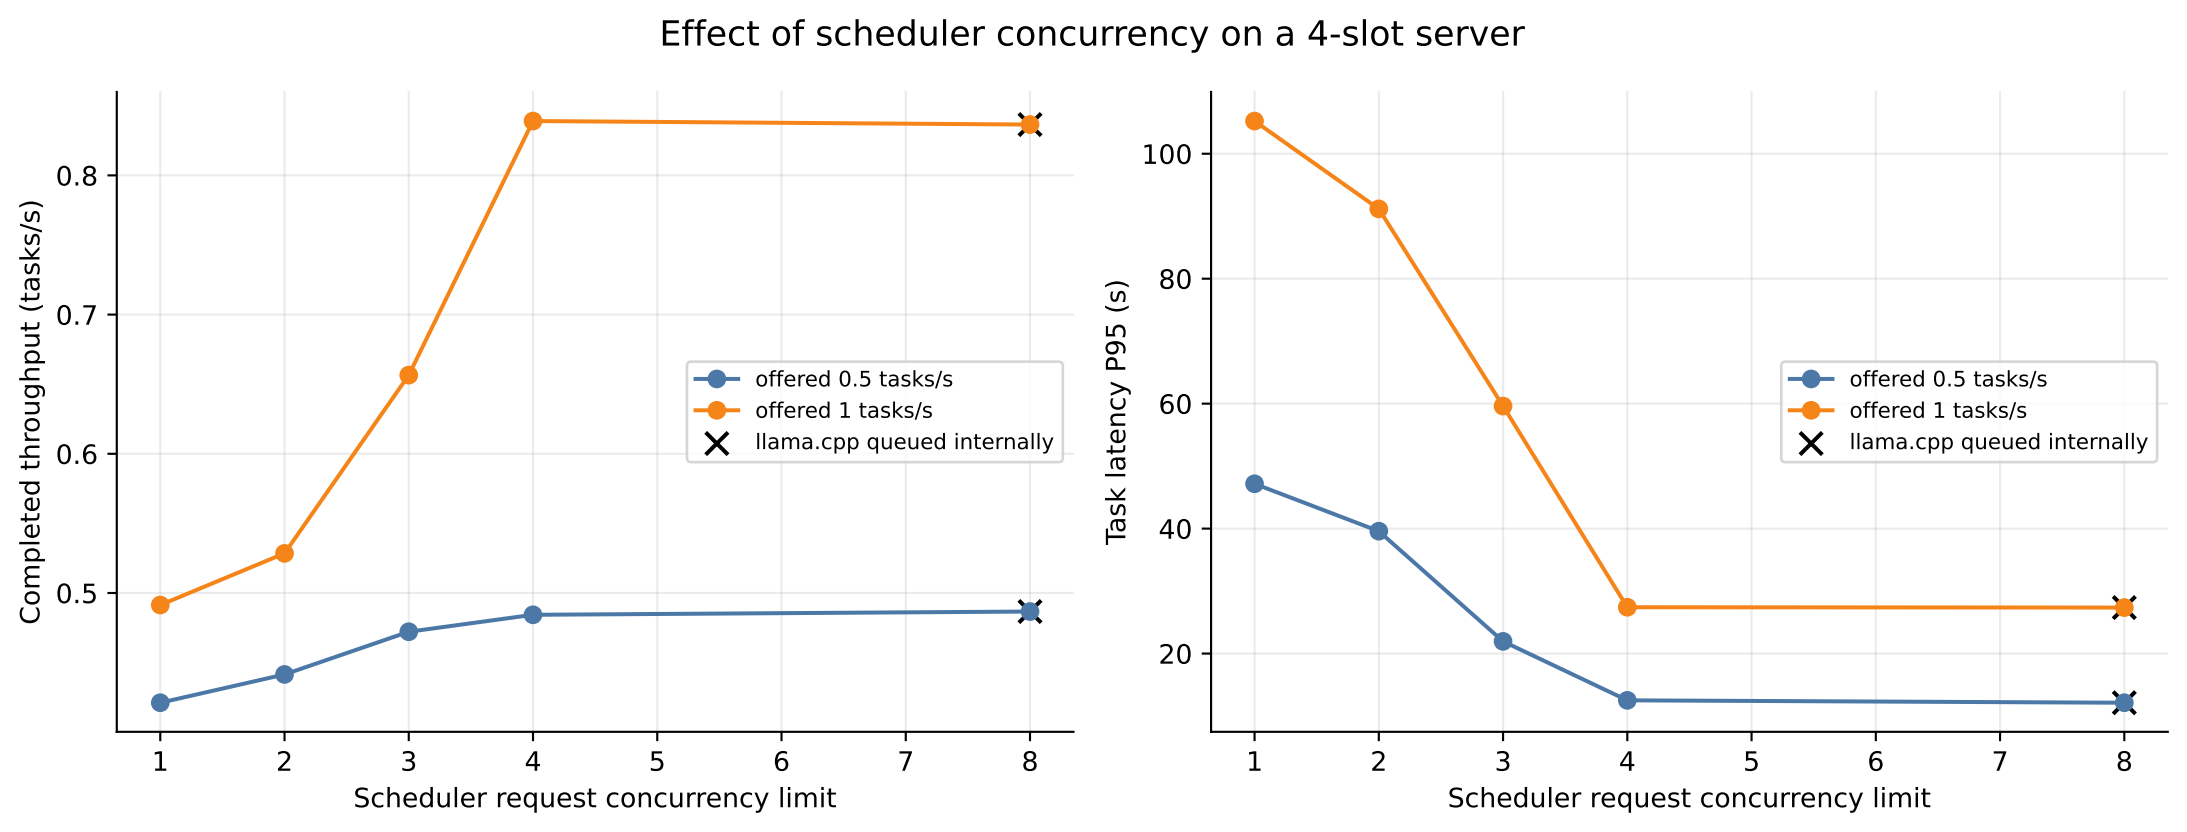
Effect of scheduler concurrency on a 4-slot server
Completed throughput (tasks/s)
0.5
0.6
0.7
0.8
1
2
3
4
5
6
7
8
Scheduler request concurrency limit
offered 0.5 tasks/s
offered 1 tasks/s
llama.cpp queued internally
Task latency P95 (s)
20
40
60
80
100
1
2
3
4
5
6
7
8
Scheduler request concurrency limit
offered 0.5 tasks/s
offered 1 tasks/s
llama.cpp queued internally